

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
from statsmodels.graphics.api import qqplot
from sklearn.metrics import mean_squared_error as mse
%matplotlib inline
```

```
↳ /usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning
import pandas.util.testing as tm
```

```
#print(sm.datasets.sunspots.NOTE)
#dta = sm.datasets.sunspots.load_pandas().data
#dta.index = pd.Index(sm.tsa.datetools.dates_from_range('1700', '2008'))
#del dta["YEAR"]
```

```
↳ ::
```

```
Number of Observations - 309 (Annual 1700 - 2008)
Number of Variables - 1
Variable name definitions::
```

```
SUNACTIVITY - Number of sunspots for each year
```


```
The data file contains a 'YEAR' variable that is not returned by load.
```

```
dta = pd.read_csv('data.csv', index_col=0, header=0, error_bad_lines=False).tail(1500).res
```

```
dta['close'].plot()
from statsmodels.tsa.arima_model import ARIMA
from sklearn.metrics import mean_squared_error
dta.head()
```

```
↳
```

	open	high	low	close	volume
0	1467.969971	1471.770020	1442.069946	1447.160034	3452650000
1	1447.550049	1456.800049	1443.729980	1447.160034	3429500000
2	1444.010010	1444.010010	1411.189941	1411.630005	4166000000
3	1414.069946	1423.869995	1403.449951	1416.180054	4221260000
4	1415.709961	1430.280029	1388.300049	1390.189941	4705390000





```
#split into test and train
percentage = 0.6
series = dta['close'].tolist()
size = int(len(series) * 0.66)
train, test = series[0:size], series[size:len(series)]
model = ARIMA(train , order = (9,0,0))
model_fit = model.fit()
```


```
from statsmodels.tsa.stattools import acf, pacf
acf_1 = acf(series)[1:20]
plt.plot(acf_1)
test_df = pd.DataFrame([acf_1]).T
test_df.columns = ["Pandas Autocorrelation"]
test_df.index += 1
test_df.plot(kind='bar')
pacf_1 = pacf(series)[1:20]
plt.plot(pacf_1)
plt.show()
test_df = pd.DataFrame([pacf_1]).T
test_df.columns = ['Pandas Partial Autocorrelation']
test_df.index += 1
test_df.plot(kind='bar')
#from the figures we conclude that it is an AR process with a lag of 8-9
```



```
/usr/local/lib/python3.6/dist-packages/statsmodels/tsa/stattools.py:541: FutureWarning
warnings.warn(msg, FutureWarning)
```



```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8b508a7240>
```



```
from keras.models import Sequential
from keras.layers import Dense,Activation,Dropout
from sklearn import preprocessing
from keras.wrappers.scikit_learn import KerasRegressor
```

```
↳ Using TensorFlow backend.
```

```
"""
```

```
Arima Rolling Forecast
```

```

"""
predicted1, resid_test = [], []
history = train
for t in range(len(test)):
    model = ARIMA(history, order=(9,0,0))
    model_fit = model.fit(dispatch=0)
    output = model_fit.forecast()
    yhat = output[0]
    resid_test.append(test[t] - output[0])
    predicted1.append(yhat)
    obs = test[t]
    history.append(obs)
    print('predicted=%f, expected=%f' % (yhat, obs))
test_resid = []
for i in resid_test:
    test_resid.append(i[0])
error = mean_squared_error(test, predicted1)
print('Test MSE: %.3f' % error)
plt.plot(test)
plt.plot(predicted1, color='red')
plt.show()

"""
Residual Diagnostics
"""
train, test = series[0:size], series[size:len(series)]
model = ARIMA(train, order=(9,0,0))
model_fit = model.fit(dispatch=0)
print(model_fit.summary())
# plot residual errors
residuals = pd.DataFrame(model_fit.resid)
residuals.plot()
plt.show()
residuals.plot(kind='kde')
plt.show()
print(residuals.describe())
#plot the acf for the residuals
acf_1 = acf(model_fit.resid)[1:20]
plt.plot(acf_1)
test_df = pd.DataFrame([acf_1]).T
test_df.columns = ["Pandas Autocorrelation"]
test_df.index += 1
test_df.plot(kind='bar')
#from the acf obtained from the residuals we conclude that
#there is still a nonlinear relationship among the residuals

```



## ARMA Model Results

```

=====
Dep. Variable:          y      No. Observations:          990
Model:                  ARMA(9, 0)    Log Likelihood      -4304.444
Method:                 css-mle      S.D. of innovations    18.665
Date:                   Wed, 15 Apr 2020    AIC                8630.887
Time:                   06:35:22    BIC                8684.762
Sample:                 0      HQIC                8651.374
=====

```

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const      1213.6978      113.115      10.730      0.000      991.997      1435.399
ar.L1.y       0.8682       0.032      27.362      0.000       0.806       0.930
ar.L2.y       0.0570       0.042       1.356      0.175      -0.025       0.139
ar.L3.y       0.0745       0.042       1.767      0.078      -0.008       0.157
ar.L4.y      -0.0045       0.042      -0.107      0.915      -0.087       0.078
ar.L5.y      -0.0412       0.042      -0.974      0.330      -0.124       0.042
ar.L6.y       0.0477       0.042       1.127      0.260      -0.035       0.131
ar.L7.y      -0.0296       0.042      -0.700      0.484      -0.113       0.053
ar.L8.y       0.0646       0.042       1.527      0.127      -0.018       0.148
ar.L9.y      -0.0412       0.032      -1.287      0.198      -0.104       0.022
=====



```


## Roots

```

=====
              Real          Imaginary      Modulus      Frequency
-----
AR.1          -1.3704          -0.0000j      1.3704      -0.5000
AR.2          -0.9471          -1.0833j      1.4389      -0.3643
AR.3          -0.9471          +1.0833j      1.4389       0.3643
AR.4          -0.0276          -1.4772j      1.4774      -0.2530
AR.5          -0.0276          +1.4772j      1.4774       0.2530
AR.6           1.0036          -0.0000j      1.0036      -0.0000
AR.7           1.0978          -1.0517j      1.5203      -0.1216
AR.8           1.0978          +1.0517j      1.5203       0.1216
AR.9           1.6898          -0.0000j      1.6898      -0.0000
=====

```





```



count  990.000000
mean    -0.327643
std     20.079317
min    -103.450298
25%     -9.477502
50%      1.216403
75%      8.601497
max     233.462208

```

```

/usr/local/lib/python3.6/dist-packages/statsmodels/tsa/stattools.py:541: FutureWarning:
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7f8acd141c88>

```



```

"""

```

```

Hybrid Model

```

```

"""

```

```

window_size = 50

```

```

def make_model(window_size):

```

```

    model = Sequential()

```

```

    model.add(Dense(50, input_dim=window_size, init="uniform",
        activation="tanh"))

```

```

    model.add(Dense(25, init="uniform", activation="tanh"))

```

```

    model.add(Dense(1))

```

```

    model.add(Activation("linear"))

```

```

    model.compile(loss='mean_squared_error', optimizer='adam')

```

```

return model

model = make_model(50)

min_max_scaler = preprocessing.MinMaxScaler()
train = np.array(train).reshape(-1,1)

train_scaled = min_max_scaler.fit_transform(test_data)


train_X,train_Y = [],[]
for i in range(0 , len(train_scaled) - window_size):
    train_X.append(train_scaled[i:i+window_size])
    train_Y.append(train_scaled[i+window_size])

new_train_X,new_train_Y = [],[]
for i in train_X:
    new_train_X.append(i.reshape(-1))
for i in train_Y:
    new_train_Y.append(i.reshape(-1))
new_train_X= np.array(new_train_X)
new_train_Y = np.array(new_train_Y)
#new_train_X = np.reshape(new_train_X, (new_train_X.shape[0], new_train_X.shape[1], 1))
model.fit(new_train_X,new_train_Y, nb_epoch=500, batch_size=512, validation_split = .05)

test_extended = train.tolist()[-1*window_size:] + test_resid
test_data = []
for i in test_extended:
    try:
        test_data.append(i[0])
    except:
        test_data.append(i)
test_data = np.array(test_data).reshape(-1,1)
min_max_scaler = preprocessing.MinMaxScaler()
test_scaled = min_max_scaler.fit_transform(test_data)
test_X,test_Y = [],[]
for i in range(0 , len(test_scaled) - window_size):
    test_X.append(test_scaled[i:i+window_size])
    test_Y.append(test_scaled[i+window_size])
    new_test_X,new_test_Y = [],[]
for i in test_X:
    new_test_X.append(i.reshape(-1))
for i in test_Y:
    new_test_Y.append(i.reshape(-1))
new_test_X = np.array(new_test_X)
new_test_Y = np.array(new_test_Y)
#new_test_X = np.reshape(new_test_X, (new_test_X.shape[0], new_test_X.shape[1], 1))
predictions = model.predict(new_train_X)
predictions_rescaled=min_max_scaler.inverse_transform(predictions)
Y = pd.DataFrame(new_train_Y)
pred = pd.DataFrame(predictions)
plt.plot(Y)
plt.plot(pred , color = 'r')

```

```
#p.plot()
plt.show()
error = mse(test_resid,predictions)
print('Test MSE: %.3f' % error)
```




Test MSE: 132.290

```
pred_final = predictions_rescaled + predicted1
error = mse(test,pred_final)
print('Test MSE: %.3f' % error)
```

Test MSE: 118.799

```
Y = pd.DataFrame(test)
pred = pd.DataFrame(pred_final)
plt.plot(Y)
plt.plot(pred , color = 'r')
#p.plot()
plt.show()
```



```
from sklearn.metrics import mean_absolute_error as mae
from sklearn.utils import check_array
#from sklearn.metrics import mean_absolute_percentage_error as mape
from sklearn.metrics import mean_squared_error
from math import sqrt
```



```
error = mse(test,pred_final)
print('Test MSE: %.3f' % error)
```

☐➤ Test MSE: 118.799

```
error = mae(test,pred_final)
print('Test MAE: %.3f' % error)
```

☐➤ Test MAE: 8.106

```
def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

```
rms = sqrt(mean_squared_error(test, pred_final))
print('Test RMS: %.3f' % rms)
```

☐➤ Test RMS: 10.900

```
print('Test MAPE: %.3f' % mean_absolute_percentage_error(test, pred_final))
```

☐➤ Test MAPE: 11.638

```
plt.plot(pred_final)
```

☐➤ [`<matplotlib.lines.Line2D at 0x7f8acc09aeb8>`]

